

Danny_Xia_HW5

February 9, 2025

```
[69]: import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import KFold
from sklearn.model_selection import cross_val_score
from sklearn.feature_selection import SequentialFeatureSelector
import matplotlib.pyplot as plt
```

```
[70]: df = pd.read_csv("heart.csv")
display(df)
```

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	\
0	63	1	3	145	233	1	0	150	0	2.3	
1	37	1	2	130	250	0	1	187	0	3.5	
2	41	0	1	130	204	0	0	172	0	1.4	
3	56	1	1	120	236	0	1	178	0	0.8	
4	57	0	0	120	354	0	1	163	1	0.6	
..	
298	57	0	0	140	241	0	1	123	1	0.2	
299	45	1	3	110	264	0	1	132	0	1.2	
300	68	1	0	144	193	1	1	141	0	3.4	
301	57	1	0	130	131	0	1	115	1	1.2	
302	57	0	1	130	236	0	0	174	0	0.0	

	slope	ca	thal	target
0	0	0	1	1
1	0	0	2	1
2	2	0	2	1
3	2	0	2	1
4	2	0	2	1
..
298	1	0	3	0
299	1	0	3	0
300	1	2	3	0
301	1	1	3	0
302	1	1	2	0

[303 rows x 14 columns]

```
[71]: X=df.drop('target', axis=1)
      y=df['target']

      X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=103/303,
      ↪random_state=42)

[72]: def ErrorEstimate(X_train_selected,X_test_selected, y_train, y_test):
      model = LogisticRegression(max_iter=3000,penalty = None)
      model.fit(X_train_selected, y_train)

      kf= KFold(n_splits=5, shuffle=True)
      cv= cross_val_score(model, X, y, cv=kf)
      print(f'Error From K-Fold Cross-Validation with {k} features::')
      print(1 - np.mean(cv))
      error = 1 - np.mean(cv)
      accuracy = model.score(X_test_selected, y_test)
      print(f"Accuracy: {accuracy}")
      test_error = 1 - accuracy
      return error, test_error

[95]: model = LogisticRegression(max_iter=3000, penalty='l1', solver='liblinear')
      model.fit(X_train, y_train)

      coefs = model.coef_[0]
      abs_coefs = np.abs(coefs)

      cv_errors = []
      test_errors = []

      for k in range(0,13):
          indices = np.argpartition(abs_coefs, -k)[-k:]

          X_train_selected = X_train.iloc[:, indices]
          X_test_selected = X_test.iloc[:, indices]

          if k == 2:
              best_columns = X_train.columns[indices].tolist()

          cross_error, test_error = ErrorEstimate(X_train_selected, X_test_selected,
          ↪y_train, y_test)
          cv_errors.append(cross_error)
          test_errors.append(test_error)

      features=list(range(1,14))
```

```

plt.figure
plt.plot(features, cv_errors, marker='o', linestyle='-', color='r',
         label='5-Fold Cross Validation Error')
plt.plot(features, test_errors, marker='o', linestyle='-', color='b',
         label='Test Error')
plt.xlabel('K')
plt.ylabel('Error')
plt.title('Error Estimation Per value of K')
plt.legend()
plt.grid(True)
plt.show()

```

```

Estimated Error From 5-Fold Cross-Validation with 0 features::
0.17158469945355181
Accuracy with 0 features: 0.8155339805825242
Estimated Error From 5-Fold Cross-Validation with 1 features::
0.171639344262295
Accuracy with 1 features: 0.6990291262135923
Estimated Error From 5-Fold Cross-Validation with 2 features::
0.17491803278688534
Accuracy with 2 features: 0.6990291262135923
Estimated Error From 5-Fold Cross-Validation with 3 features::
0.18513661202185783
Accuracy with 3 features: 0.7184466019417476
Estimated Error From 5-Fold Cross-Validation with 4 features::
0.1750273224043717
Accuracy with 4 features: 0.7281553398058253
Estimated Error From 5-Fold Cross-Validation with 5 features::
0.17158469945355193
Accuracy with 5 features: 0.7572815533980582
Estimated Error From 5-Fold Cross-Validation with 6 features::
0.17486338797814205
Accuracy with 6 features: 0.7864077669902912
Estimated Error From 5-Fold Cross-Validation with 7 features::
0.16169398907103827
Accuracy with 7 features: 0.8155339805825242
Estimated Error From 5-Fold Cross-Validation with 8 features::
0.1654098360655738
Accuracy with 8 features: 0.7961165048543689
Estimated Error From 5-Fold Cross-Validation with 9 features::
0.1618579234972678
Accuracy with 9 features: 0.7864077669902912
Estimated Error From 5-Fold Cross-Validation with 10 features::
0.1680874316939891
Accuracy with 10 features: 0.8058252427184466
Estimated Error From 5-Fold Cross-Validation with 11 features::

```

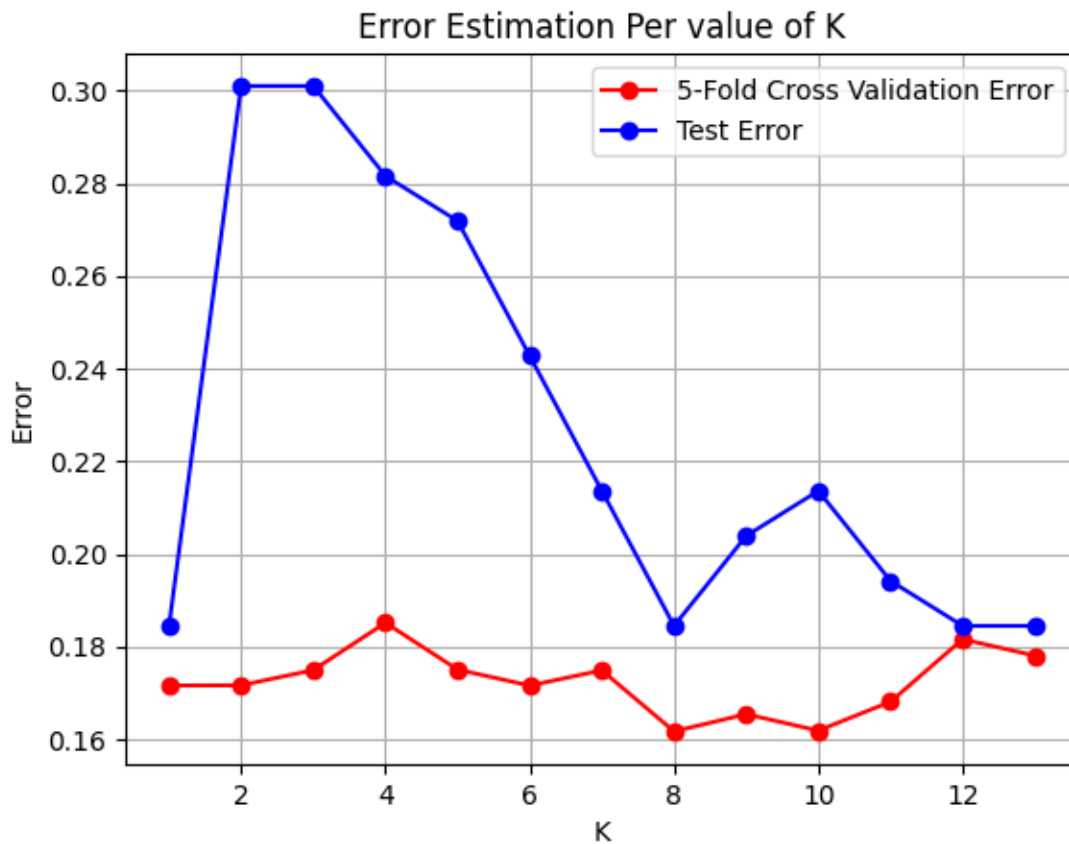
0.18158469945355193

Accuracy with 11 features: 0.8155339805825242

Estimated Error From 5-Fold Cross-Validation with 12 features::

0.17803278688524604

Accuracy with 12 features: 0.8155339805825242



```
[98]: from matplotlib.colors import ListedColormap

def plot_decision_boundary(X, y, model, selected_features):
    X_selected = X[selected_features]
    x_min, x_max = X_selected.iloc[:, 0].min() - 1, X_selected.iloc[:, 0].max()
    ↪+ 1
    y_min, y_max = X_selected.iloc[:, 1].min() - 1, X_selected.iloc[:, 1].max()
    ↪+ 1
    xx, yy = np.meshgrid(np.linspace(x_min, x_max, 100), np.linspace(y_min,
    ↪y_max, 100))

    Z = model.predict(np.c_[xx.ravel(), yy.ravel()])
    Z = Z.reshape(xx.shape)
```

```

plt.contourf(xx, yy, Z, alpha=0.3, cmap=ListedColormap(['#FFAAAA',
↪ '#AAAAFF']))
plt.scatter(X_selected.iloc[:, 0], X_selected.iloc[:, 1], c=y,
↪ edgecolors='k', cmap=ListedColormap(['#FF0000', '#0000FF']))
plt.xlabel(selected_features[0])
plt.ylabel(selected_features[1])
plt.title("Decision Boundary for k=2")
plt.show()


print(best_columns)
model = LogisticRegression(max_iter=1500, penalty=None)
model.fit(X_train[best_columns], y_train)

plot_decision_boundary(X_train, y_train, model, best_columns)

```

```
['thal', 'ca']
```

/Library/Frameworks/Python.framework/Versions/3.12/lib/python3.12/site-packages/sklearn/utils/validation.py:2739: UserWarning: X does not have valid feature names, but LogisticRegression was fitted with feature names
warnings.warn(



[]: